


# The Circumstantial Event Ontology (CEO.owl) and ECB+/CEO Evaluation Corpus

Roxane Segers\*, Tommaso Caselli\*\*, Piek Vossen\*

\*Vrije Universiteit Amsterdam, The Netherlands, \*\*Rijksuniversiteit Groningen, The Netherlands



We do  
semantic  
event  
domino!

"Three months ago, Mary B. **arsoned** the cottage to punish her ex-husband."  
"The insurance company made sure the house was **restored** by an expert builder."

So, why was the house restored?

## What it is:

CEO.owl is an event ontology that:

- a) types events and identifies the roles of the entities related to the event;
- b) models what situation is true before, during and after an event;
- c) allows for chaining event mentions with semantic circumstantial relations, based on shared properties in the situation assertions.

## What it does:

The model allows to infer semantic circumstantial relation chains between event mentions in natural language text.

## What it needs:

- Minimally: the vocabulary mappings and basic chaining heuristics;
- For full expressivity: FrameNet based SRL output and a Reasoner.

## Domain:

Calamities in newswire: murders, earthquakes, theft, cyber attacks, etc.

## What is a circumstantial event relation?

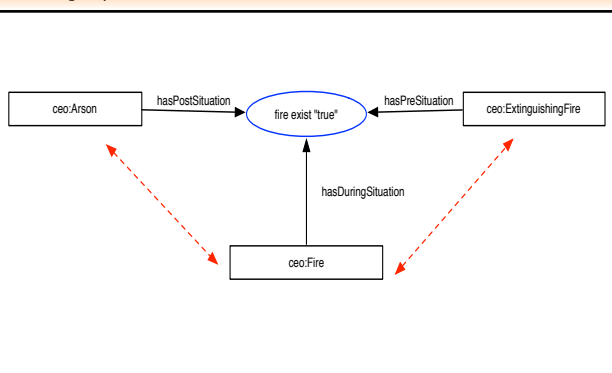
A weak causal relation. Event A gives rise to event B, but not necessarily so:

*X being hit — X being injured — X dying;  
Z damaging Y — Q rebuilding Y — Y repaired*

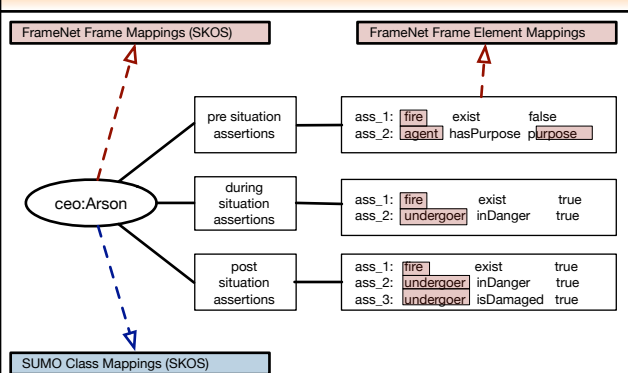
## How do we infer those relations?

WHEN two event mentions appear in the same textual context,  
we infer a circumstantial relation;  
ONLY IF the ceo event classes the mentions pertain to  
share properties in the situation assertions.

## Inferring implicit circumstantial relations



## The CEO metamodel and mappings



Available: <http://github.com/newsreader/eso-and-ceo>

## Publicly available (CC-BY-SA):

### CEO.owl

- 223 event classes
- 121 properties
- 79 roles
- 456 situation rule assertions
- skos mappings to SUMO (195) and FrameNet Frames (265) and Frame Elements (263)

### ECB+/CEO, the manually annotated evaluation corpus

- 508 articles
- 2494 CEO links annotated, related to ± 9500 event mentions
- 3448 new coreference sets

### CEO vocabulary (version 1)

- 650 unique surface forms expressing events mapped to CEO.

## First evaluation results

### Evaluation 1: based on event mentions

Baseline:  
P 0.28  
R 0.99  
F1 0.43

System:  
P 0.42  
R 0.11  
F1 0.18

### Evaluation 2: based on event instances

Baseline:  
P 0.28  
R 0.99  
F1 0.43

System:  
P 0.64  
R 0.24  
F1 0.24

recall: - 35% of annotated relations is NOT circumstantial but episodic  
- additional heuristics needed to connect mention to mention B  
- extension needed of ceo event assertions

precision: - annotators missed relations, but correctly retrieved by the system (10%)  
- vocabulary mapping issues: polysemy